

Ian Piper  
CSCI203

# Algorithms and Data Structures

---

Week 3 – Lecture A

# Applied Algorithms

---

- ▶ Simulation

- ▶ Two basic kinds:

- ▶ Continuous (time steps by a constant amount)
    - ▶ Discrete (time steps from event to event)

- ▶ Each type is suited to different kinds of simulations

- ▶ Continuous – processes where there are no time-based events

- ▶ Physical processes (e.g. Explosions)

- ▶ Discrete – processes where time-based events control what is happening.

- ▶ (e.g. Queues)



## Note

---

- ▶ This terminology is not especially obvious or clear.
- ▶ Continuous:
  - ▶ Time is broken into discrete chunks (ticks)
  - ▶ Usually used to model a continuous process.
- ▶ Discrete:
  - ▶ Time can take any value
  - ▶ Usually used to model discrete events

# Continuous Simulation

---

Clock Driven



## Continuous Simulation.

---

- ▶ Often dependent on complex mathematics.
- ▶ Often based on gridded algorithms.
  - ▶ Implicit solution
  - ▶ Explicit solution
- ▶ Often requires extreme computing resources.
  - ▶ Supercomputers
- ▶ We will not be looking at continuous simulation in this subject.

## Examples of Continuous Simulation: 1

---

- ▶ Mine explosions
  - ▶ Solve 10 differential equations for each of several thousand cells for every millisecond of the simulated event.
  - ▶ Parallel supercomputer.
  - ▶ Still over 24 hours per run.
- ▶ Some videos of the results...



## Examples of Continuous Simulation: 2

---

- ▶ Social simulation
  - ▶ Track the state of individuals in a simulated environment.
- ▶ Two entities
  - ▶ Peeps – simulated individuals (not always people)
  - ▶ Cells – simulated locations
- ▶ Used to examine response to threat
  - ▶ Spread of disease
  - ▶ Natural disaster
  - ▶ Man-made disaster (terrorist attack)
- ▶ Again, lots of time/computer power required.

# Discrete Simulation

---

Event driven



# Discrete Simulation

---

- ▶ Normally a lot less mathematically complex.
- ▶ Usually requires a lot less computer resources.
- ▶ E.g. Queue simulation
  - ▶ Widely used to evaluate queue-based processes
    - ▶ Shops
    - ▶ Production lines
    - ▶ Industrial processes
  - ▶ We will look at some simple examples and see how we might implement them.

## Scenario 1:

---

- ▶ A single server queue.
  - ▶ Customers arrive at random intervals to be served
  - ▶ If the server is not busy the customer will be served immediately
  - ▶ If the server is busy the customer will join the end of the (possibly empty) queue.
  - ▶ When the server has finished with the customer the next customer (if any) begins service – first customer in queue.



## Scenario 1:

---

- ▶ Events.
  - ▶ Customer arrives
  - ▶ Customer starts service
  - ▶ Customer ends service and leaves
- ▶ What we know:
  - ▶ When each customer arrives
  - ▶ How long they take to serve
- ▶ What we want to know:
  - ▶ How big is the queue on average?
  - ▶ How busy is the server?
  - ▶ What proportion of customers have to wait in a queue?

## Scenario 1:

---

- ▶ Data

- ▶ Input data is a file consisting of a set of records containing

- ▶ Arrival time

- ▶ Service time

- ▶ For each customer

- ▶ Sorted by arrival time

- 0.24 0.55

- 0.59 0.16

- 0.90 0.07

- 1.87 0.69

- ...



## Scenario 1:

---

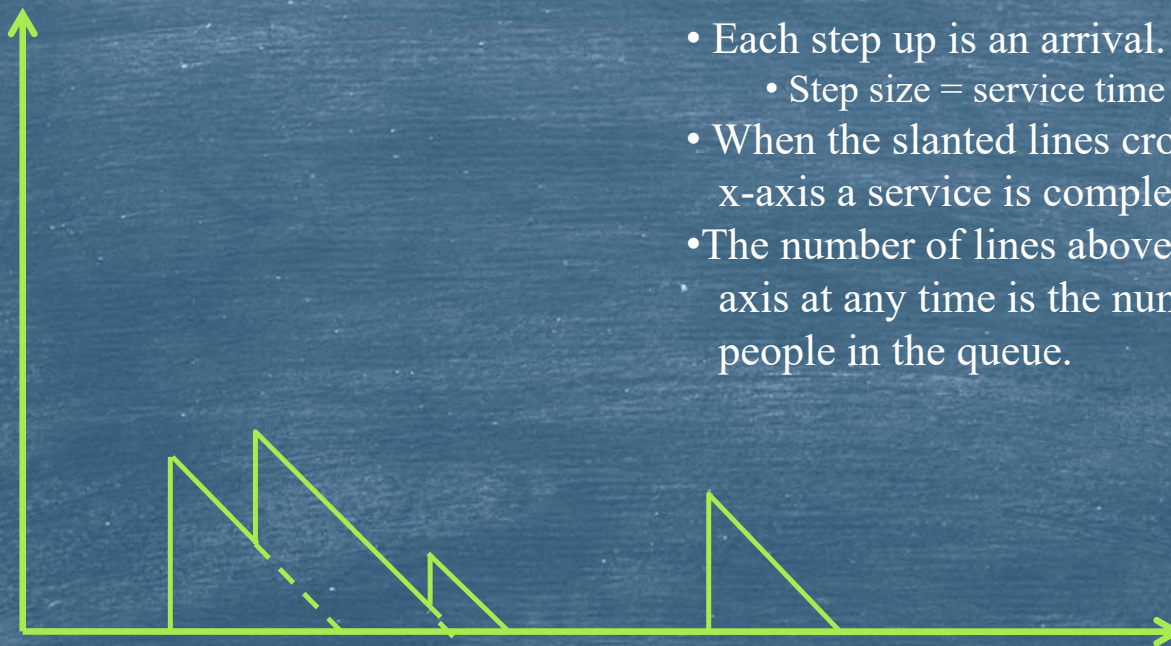
### ► Manual Simulation

- From the data file we can get a feel for what is happening
- At time 0.00 the simulation starts
  - The server is idle
  - The queue is empty
- At time 0.24 the first customer arrives
  - The server is busy for the next 0.55 (until 0.79)
  - The queue is empty
- At time 0.59 the second customer arrives
  - The server is still busy
  - The queue now contains 1 customer (customer 2)
- At time 0.79 the server finishes with customer 1
  - The server stays busy for the next 0.16 (until 0.95)
  - The queue is empty

0.24	0.55
0.59	0.16
0.90	0.07
1.87	0.69
...	

# Graphs of Queues

- We can represent what is happening with a graph:



- Each step up is an arrival.
  - Step size = service time
- When the slanted lines cross the x-axis a service is complete.
- The number of lines above the axis at any time is the number of people in the queue.



# The Algorithm

---

- ▶ Starting to design the algorithm.
  - ▶ Data structures
    - ▶ We need to hold the queue
      - ▶ What should we put in it?
    - ▶ We need to keep track of the time
    - ▶ We need to know if the server is busy
      - ▶ If so we need to know when they will finish
    - ▶ We need to know when the next customer will arrive
    - ▶ We need to track statistics

# The Algorithm

---

- ▶ Starting to design the algorithm.
  - ▶ Procedures
    - ▶ Initialise the simulation
    - ▶ Process an arrival
    - ▶ Process a service completion
    - ▶ Finish the simulation
- ▶ Once the simulation is running how do we decide what to do next?
  - ▶ Compare the next arrival time with the end of service time



# The Algorithm

---

- ▶ Initialise the simulation

- time = 0
  - busy = false
  - queue = empty
  - read next\_arrival, next\_service

- ▶ Set up for statistics collection

# The Algorithm

---

## ► The main program loop

```
initialise
repeat
    if busy = true then
        if service_end < next arrival then
            process_service_end
        else
            process_arrival
    else
        process_arrival
    fi
until arrival file is empty and busy = false
finish
```



# The Algorithm

---

## ► Process an arrival

```
time = next_arrival
if busy then
    enqueue (next_service)
else
    busy = true
    service_end = time + next_service
fi
read (next_arrival next_service)
```

# The Algorithm

---

## ► Process a service completion

```
time = service_end
if queue_empty then
    busy = false
else
    service_end = time + dequeue()
fi
```



# Getting Complicated

---

- ▶ What if there is more than one server?
  - ▶ Two possible situations
    - One queue for all servers (like a bank)
    - One queue per server (like a supermarket)

## "One Queue to Serve Them All..."

---

- ▶ One queue for all servers
  - ▶ If all servers are busy add arrival to queue
    - ▶ Otherwise make one of the idle servers busy
  - ▶ When a server finishes have them serve the head of the queue
    - ▶ Otherwise make them idle



# Single Queue

---

- ▶ The events we are interested in are now:
  - ▶ Customer arrives
  - ▶ Server 1 finishes
  - ▶ Server 2 finishes
  - ▶ ...
  - ▶ Server  $n$  finishes
- ▶ How do we keep track of which event will happen next?
- ▶ Do we really need to know which servers are busy or can we just keep track of how many are busy?

# Single Queue

---

## ► Keeping track of servers:

### ► Two possible solutions

- An array of servers with `busy[i]` and `end_time[i]`  
This lets us track who is doing what  
Finding what happens next is in  $O(n)$
- A heap of end times (smallest on top)  
This does not let us track who is doing what  
Finding what happens next is  $O(\log n)$
- Can we get the best of both worlds?



# Single Queue

---

- ▶ Using a heap
  - ▶ If we are clever, we can store all event times on the heap
  - ▶ All we need is a second array telling us what is what
  - ▶ If we are really clever, we can partition the heap into two parts:
    - The heap itself
    - The idle servers
  - ▶ How do we manipulate the heap as events occur?

# Single Queue

---

## ► Customer arrives:

```
time = heap[0]
read next arrival into heap[0], next_service_time
sift_down(heap)
if n_busy < n_servers then
    n_busy = n_busy + 1
    heap[n_busy] = time + service_time
    sift_up(heap)
else
    enqueue(service_time)
fi
service_time = next_service_time
```



# Single Queue

---

► Server finishes:

```
time = heap[0]
if queue_empty then
    heap[0] = heap[n_busy]
    n_busy = n_busy - 1
else
    heap[0] = time + dequeue()
fi
sift_down(heap)
```

# Single Queue

---

- ▶ In summary:
  - ▶ Heap Grows if a customer arrives and a server is idle
  - ▶ Heap shrinks if a customer is served and the queue is empty
  - ▶ Heap stays the same size otherwise
- ▶ If we keep a second array (id) initially filled with integers  $0..n$  we can use it to track who is doing what
  - ▶ 0 is the next arrival
  - ▶ 1 is server 1's completion time
  - ▶ 2 is server 2's completion time
  - ▶ ...
  - ▶  $n$  is server  $n$ 's completion time



# Single Queue

---

## ► Customer arrives:

```
time = heap[0]
read next arrival into heap[0], next_service_time
sift_down(heap, id)
if n_busy < n_servers then
    n_busy = n_busy + 1
    heap[n_busy] = time + service_time
    sift_up(heap, id)
    service_time = next_service_time
else
    enqueue(next_service_time)
fi
```

# Single Queue

---

## ► Server finishes:

```
time = heap[0]
if queue_empty then
    swap (id[0], id[n_busy])
    heap[0] = heap[n_busy]
    n_busy = n_busy - 1
else
    heap[0] = time + dequeue()
fi
sift_down(heap, id)
```



## Single Queue

---

- ▶ Every time we move an entry in the heap we move the corresponding entry in the id array.
- ▶ If the top of the id array is a zero the next event is an arrival
- ▶ If the top of the id array is non\_zero the next event is a service completion for server id[0]
- ▶ The simulation starts with the first arrival time in heap[0] and n\_busy = 0

## Multiple Queues

---

- ▶ In this case we have an array of  $n$  queues, 1 per server
- ▶ When a customer arrives and all servers are busy we place the customer on one of the queues (which one?)
- ▶ When a server finishes we only make them busy if their queue is not empty
- ▶ NOTE: This means that we can have a queue even if a server is idle.



# Priority queues

---

- ▶ How would we handle the situation where customers are given different service priorities?
  - One queue for each priority empty the highest priority queue first
  - This is only efficient if there are a small number of priorities
- ▶ What do we do if each priority may be different?
  - E.g. priority is a float between 0 and 1
  - 0 is the lowest customer priority
  - 1 is the highest priority customer
  - We have an infinite number of different priorities so we can't have a queue for each one.

## Priority queues

---

- ▶ The solution is to replace the queue with a heap ordered on priority
- ▶ Each time we remove a customer from the heap we
  - ▶ Move the last entry to the top of the heap
  - ▶ Reduce the heap size by one
  - ▶ Sift down the top entry
- ▶ Each time we add a customer to the heap we:
  - ▶ Increase the heap size by one
  - ▶ Add the customer to the end of the heap
  - ▶ Sift up the last entry