


Ian Piper
CSCI203

Algorithms and Data Structures

Week 4 – Lecture A

Binary Trees Revisited

- ▶ Remember from 1st year.
- ▶ A Binary tree is a tree in which each node has a maximum of two child nodes, the left child and the right child.



Binary Trees

- ▶ A binary tree can be implemented in several ways.

- ▶ Some useful approaches are:

- ▶ In an array:

- ▶ `tree: array of stuff`

- ▶ Root is `tree[1]`

- ▶ The children of `tree[i]` are `tree[2*i]` and `tree[2*i+1]`

- ▶ As a collection of dynamic records:

- ▶ `type tree_node = record`

- `contents: stuff`

- `left: ^tree_node`

- `right: ^tree_node`

- ▶ `root: ^tree_node`

- ▶ As an array of records:

- ▶ `type tree_array_node = record`

- `contents: stuff`

- `left: int`

- `right: int`


- ▶ `tree: array of tree_array_node`

Binary Search Trees

Binary Search Tree

- ▶ This is a binary tree with one extra condition:
 - ▶ For each non-leaf node:
 - ▶ The contents of the left child \leq the contents of the node;
 - ▶ The contents of the node \leq the contents of the right node.

▶ E.g.




Searching a BST

```
► procedure find(value: stuff, node: ^tree_node): ^tree_node
    if value == nil then
        return not_found
    fi
    if value == node.contents then
        return node
    else if value < node.contents then
        find(value, node.left)
    else
        find(value, node.right)
    fi
end
```


Searching: An Example

- ▶ Consider the BST shown at the right:
- ▶ `find(13, root)`
- ▶ `13 < 15; find(13, root.left)`
- ▶ `find(13, node)`
- ▶ `13 > 9; find(13, node.right)`
- ▶ `find(13, node)`
- ▶ `13 = 13; return node`



Building a BST

- ▶ To build a BST we add nodes, one at a time by:
 - ▶ Searching the existing BST for the value to be inserted.
 - ▶ If, the value is not found:
 - ▶ Create a new node;
 - ▶ Add the new node to the tree as the appropriate child of the last node examined.
- ▶ A comparison between the value to be inserted and the value stored in the last valid node will determine which child is to be selected.
- ▶ The first node is a special case:
 - ▶ Here we must create the first node of the tree and point root at it.

Building a BST

```
► procedure insert(value: stuff, node): ^tree_node
    next: ^tree_node, left: boolean
    if value == node.contents then
        return // already in the tree
    else if value < node.contents then
        next = node.left; left = true // we need to go left
    else
        next = node.right; left = false // we need to go right
    fi
    if next != nil then
        insert (value, next) // keep trying
    else
        next = new_tree_node // make a new node
        next.contents = value // store the value
        if left then // update the parent
            node.left = next
        else
            node.right = next
        fi
    fi
fi
end
```

Building a BST

```
▶ procedure insert_first(value): ^tree_node
    node: ^tree_node
    start = new_tree_node
    start.contents = value
    return start
end
```

- ▶ We create the BST by:
 - ▶ `root = insert_first(value)`

Building: An Example

- ▶ Let us build a BST from the following values:

- ▶ 15, 33, 9, 13, 5, 21, 11

- ▶ `root = insert_first(15)`

- ▶ `insert(33)`


- ▶ `insert(9)`

- ▶ `insert(13)`

- ▶ `insert(5)`

- ▶ `insert(21)`

- ▶ `insert(11)`



Sorting With BSTs


- ▶ If we perform an in-order traversal of a binary search tree, the nodes of the tree will be listed in sorted order.
- ▶ Recall:
 - ▶ In order traversal:
 - ▶

```
procedure visit(node: ^tree_node)
  if node.left != nil then
    visit(node.left)
  fi
  print(node.contents)
  if node.right != nil then
    visit(node.right)
  fi
  return
end
```
 - ▶ This is BST Sort – simply call **visit(root)**.

Sorting: An Example

- ▶ Consider the BST shown to the right:


```
▶ visit(root)           ▶ return
▶ visit(root.left)      ▶ return
▶ visit(node.left)      ▶ print(node.contents)
▶ print(node.contents)  ▶ visit(node.right)
▶ return               ▶ visit(node.left)
▶ print(node.contents)  ▶ print(node.contents)
▶ visit(node.right)     ▶ return
▶ visit(node.left)      ▶ print(node.contents)
▶ print(node.contents)  ▶ return
▶ return               ▶ return
▶ print(node.contents)  ▶ return
```




5	9	11	13	15	21	33
---	---	----	----	----	----	----

The Problem with BSTs

- ▶ If we create a BST from the following sequence:



- ▶ We get the following BST:



- ▶ This tree is severely unbalanced.

Balancing a BST

Balancing a BST


- ▶ Operations on BSTs are only $\Theta(\log(n))$ for balanced trees.
- ▶ Can we adjust a BST, as we operate on it, to keep it more or less balanced?
- ▶ How efficient is this?
- ▶ Is it worth the effort?
- ▶ What do we mean by balanced, anyway?

Balanced?

- ▶ Try “left and right subtrees must be of the same height”
 - ▶ This is easy to achieve but not very useful.
 - ▶ Too soft.
- ▶ Try “every node must have left and right subtrees of the same height”
 - ▶ This is impossible unless the tree is complete.
 - ▶ Too hard.
- ▶ Try “every node must have left and right subtrees which differ in height by at most 1”
 - ▶ This is the AVL (Adelson-Velski and Landis) balance condition.
 - ▶ Just right. (As we shall now see.)


AVL Trees

- ▶ This is an AVL tree:
- ▶ These are the heights:
- ▶ Note that, at each node, the heights of its children differ by at most one.



AVL Trees

- ▶ This is not an AVL tree:
- ▶ These are the heights:
- ▶ Note that the root node has children which differ in height by two.




Losing My Balance

- ▶ Insertion can unbalance an AVL tree node β
 1. An insertion into the left subtree of the left child of β
 2. An insertion into the right subtree of the left child of β
 3. An insertion into the left subtree of the right child of β
 4. An insertion into the right subtree of the right child of β
- ▶ Cases 1 and 4 are equivalent, as are cases 2 and 3 (although there are still 4 cases from a coding viewpoint).


An Abstract Tree

- ▶ Consider the following abstract tree:
- ▶ Triangles represent sub-trees:
- ▶ Boxes represent nodes.




Tree Rotation

- ▶ We define a simple operation on this tree; rotation.
- ▶ As we can see:
 - ▶ Node y replaces node x as the “local root”.
 - ▶ Node x becomes the right child of node y.
 - ▶ Subtree D moves from being the right child of node y to the left child of node x.
- ▶ This is a *right* rotation.
- ▶ The pivot node becomes the right child.
- ▶ Note: If we started with a BST we still have a BST!




Case 1


- Consider what happens when insertion into the left subtree of node y causes the tree to lose its AVL balance at node x :
- Tree heights are shown below the components:



Before




After




We have restored the balance!

Case 2


- Consider what happens when the imbalance occurs as a result of insertion into the right subtree of the left child:
- Will a single rotation work this time?



Before




After




We still have an imbalance!

Case 2

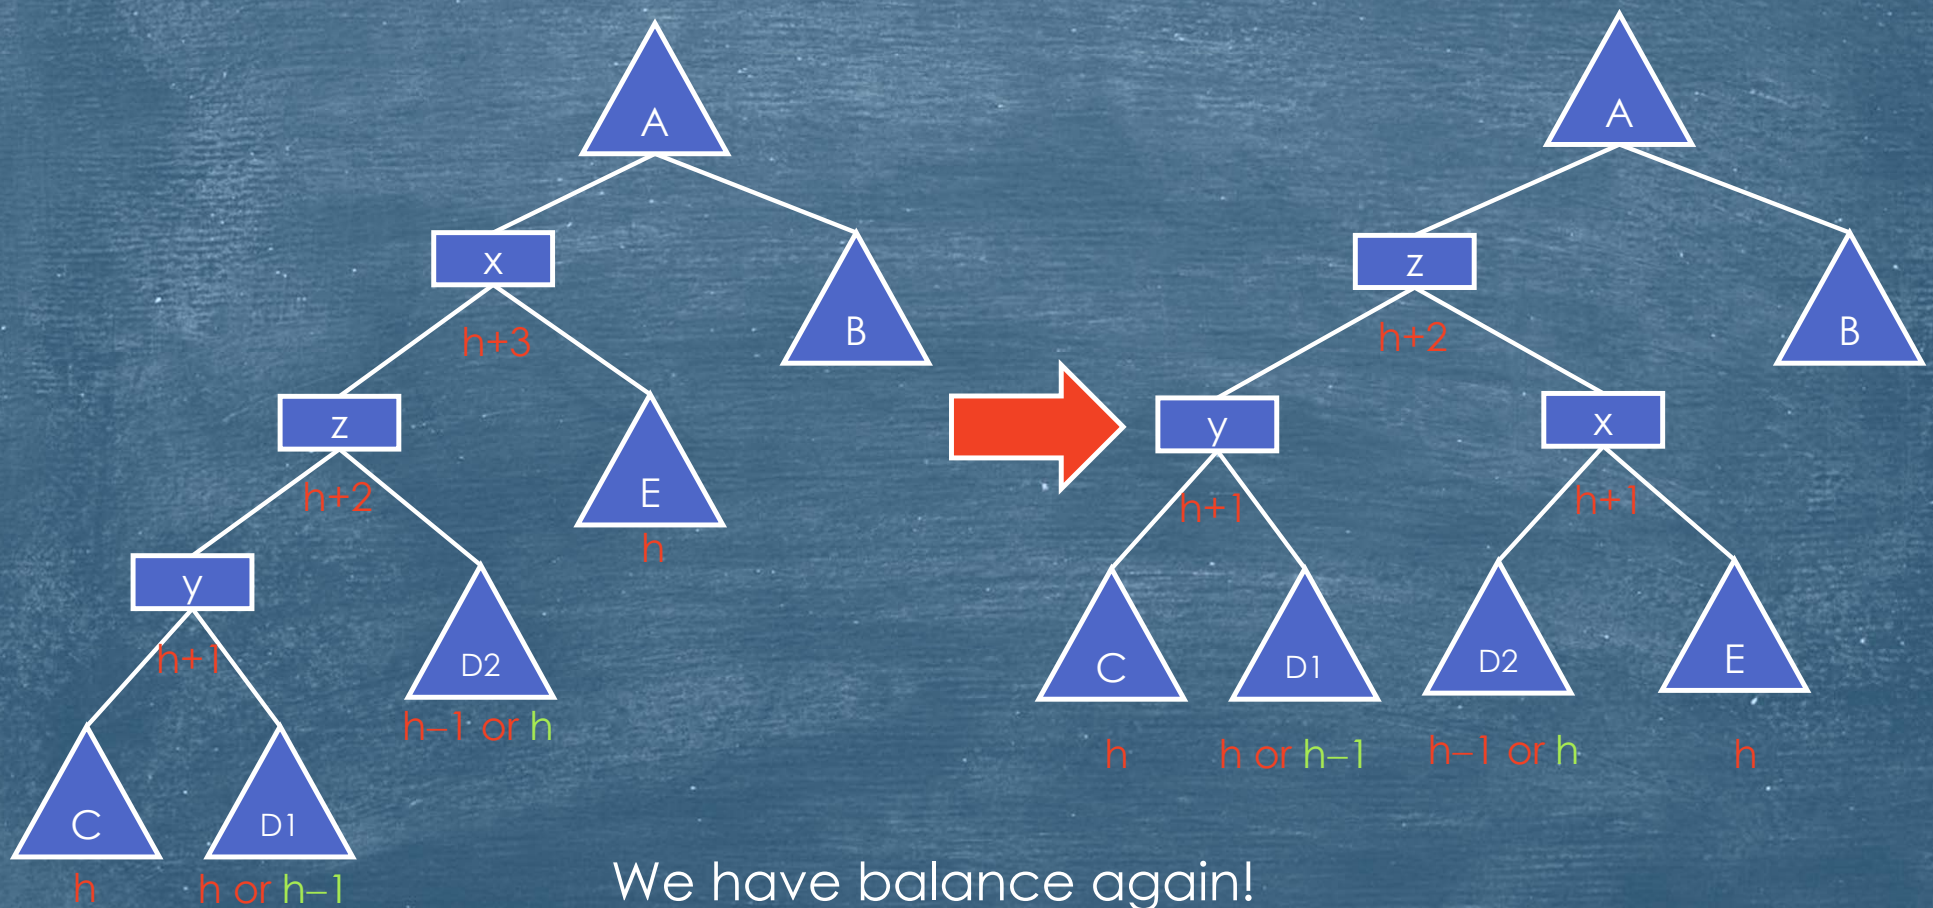
- Let us expand subtree D.



First perform a *left* rotation on node y



Then perform a *right* rotation on node x



An Example

► Let us build an AVL tree one node at a time:

► Insert 3

► Insert 2

► Insert 1

► We have an imbalance at node 3


► Right rotate node 3

► Insert 4

► Insert 5

► We have an imbalance at node 3

► Left rotate node 3



Continued


- ▶ The tree so far

- ▶ Insert 6

- ▶ We have an imbalance at node 2
 - ▶ Left rotate node 2


- ▶ Insert 7

- ▶ We have an imbalance at node 5
 - ▶ Left rotate node 5




Continued

- ▶ The tree so far
 - ▶ Insert 16
 - ▶ Insert 15
 - ▶ We have an imbalance at node 7
 - ▶ A double rotation is needed
 - ▶ First, right rotate node 16
 - ▶ Then, left rotate node 7



Continued

- ▶ The tree so far
 - ▶ Insert 14
 - ▶ We have an imbalance at node 6
 - ▶ A double rotation is needed
 - ▶ First, right rotate node 15
 - ▶ Then left rotate node 6
- ▶ And so on.



AVL Trees – Implementation

```
► type avl_node = record  
    value: stuff  
    left:  ^avl_node  
    right: ^avl_node  
    height: int
```


AVL Trees – Implementation

```
► procedure avl_insert(key, tree)
    if (tree = nil) then
        tree = new avl_node(key, nil, nil, 0)
    else if key < tree.value then
        avl_insert(key, tree.left)
        if (tree.left).height - (tree.right).height = 2 then
            if key < (tree.left).value then
                rotate_right(tree) // case 1
            else
                double_right(tree) // case 2
            fi
        else if key > tree.value then
            avl_insert(key, tree.right)
            if (tree.right).height - (tree.left).height = 2 then
                if key < (tree.right).value then
                    double_left(tree) // case 3
                else
                    rotate_left(tree) // case 4
                fi
            fi
        fi
        tree.height = max((tree.left).height, (tree.right).height) + 1
    end
```

AVL Trees – Implementation

- ▶

```
procedure rotate_right(k2)
    k1 = k2.left
    k2.left = k1.right
    k1.right = k2
    k2.height = max((k2.left).height), (k2.right).height) + 1
    k1.height = max((k1.left).height), k2.height) + 1
    k2 = k1
end
```
- ▶

```
procedure rotate_left(k2)
    k1 = k2.right
    k2.right = k1.left
    k1.left = k2
    k2.height = max((k2.left).height), (k2.right).height) + 1
    k1.height = max(k2.height, (k1.right).height), ) + 1
    k2 = k1
end
```


AVL Trees – Implementation

- ▶

```
procedure double_right( k3)
    rotate_left(k3.left)
    rotate_right(k3)
end
```

- ▶

```
procedure double_left( k3)
    rotate_right(k3.right)
    rotate_left(k3)
end
```

- ▶ Note: The pseudo-code presented here does not take into account the fact that the parent of the top node must change as part of the rotation.

- ▶ Implementation of this detail is left as an exercise. ☺