

Ian Piper
CSCI203

Algorithms and Data Structures

Week 6 – Lecture A

Hashing – Picking m

Hashing: Picking m

- ▶ As we saw last week we want m , the number of slots in the dictionary, to be $\Theta(n)$, where n is the number of entries in the dictionary.
- ▶ Remember: operations on a dictionary are $O(1+n/m)$, so if n grows too large we get less and less efficient.
- ▶ The problem we face is that, often, we do not know how many records we will need to store.
 - ▶ If m is too small, the dictionary becomes inefficient.
 - ▶ If m is too large, we waste storage (memory or disc).
- ▶ How do we get the right value for m ?
- ▶ Let's say we want $m \geq n$ at all times.

Lucky Guess?

- ▶ If we have no knowledge of the ultimate size of n , what can we do?
 - ▶ Guess.
 - ▶ Pick m based on an optimistic assessment of the likely size of n .
 - ▶ No idea?
 - ▶ Pick your favourite small number ☺.
 - ▶ $m = 8$, say.
 - ▶ Now what?
 - ▶ What if n turns out to be greater than 8?
 - ▶ Make m bigger.
 - ▶ How much bigger?

Changing m

- ▶ Hang on a sec.
- ▶ If we change m we have problems:
 - ▶ Our hash array is too small.
 - ▶ Our hashed keys will be wrong.
 - ▶ They depend on the value of m .
- ▶ Does this mean that we have to recreate the hash table from scratch?
 - ▶ It sure does.
- ▶ Isn't this a BAD THING™?

Growing a Hash Table

Growing a Hash Table

- ▶ What exactly has to happen if we change m ?
 - ▶ Let's say the new table size is m' .
- ▶ We now need a new array with m' elements.
 - ▶ We also need to move all of the existing elements from the old table to the new one.
- ▶ Build a new hash function h' .
 - ▶ Remember, the hash function depends on m .
- ▶ Insert the existing data into the new table.
 - ▶ This involves re-hashing every key.
- ▶ So, the first question is:
 - ▶ How much do we grow m ?

$$m' = ?$$

- ▶ Each time we grow the table we perform $\Theta(m+n+m')$ operations.
 - ▶ This is $\Theta(n)$.
- ▶ Let's look at some options:
 - ▶ $m' = m+1$.
 - ▶ What is the cost of n insertions?
 - ▶ $\Theta(1)$ for the first m insertions.
 - ▶ $\Theta(m')$ for each insertion after that.
 - ▶ Overall $\Theta(n^2)$

$$m' = ?$$

- ▶ $m' = 2m$
 - ▶ $\Theta(1)$ for the first m insertions.
 - ▶ $\Theta(m)$ for the next insertion.
 - ▶ $\Theta(1)$ for the next $m-1$ insertions.
 - ▶ $\Theta(2m)$ for the next insertion.
 - ▶ $\Theta(1)$ for the next $2m-1$ insertions.
- ▶ Overall $\Theta(n + (n/2) + (n/4) + \dots) = \Theta(2n) = \Theta(n)$
- ▶ The cost of expanding the table gets spread over the extra elements we are making room for.
- ▶ This is known as *Amortized* cost.
- ▶ Note: an amortized cost of $\Theta(1)$ per operation does not mean that every operation has this cost.
 - ▶ Just that this is the average cost per operation.

Amortized Cost

Amortized Cost

- ▶ We say an operation has a cost of “ $T(k)$ Amortized” if k operations take a total of $k \times T(k)$ time.
- ▶ Table doubling takes $\Theta(n)$ operations for n insertions so the amortized cost is $\Theta(1)$.
- ▶ This is, actually, a GOOD THING™.
- ▶ Note: we can use table doubling to implement any solution where we do not know the size of the data structure in advance and it grows in a “well behaved” way.
- ▶ Table doubling minimizes the cost associated with dynamic data structures.

Deletions

- ▶ What about deletions?
 - ▶ Each deletion is still $\Theta(1)$.
 - ▶ They simply increase the number of operations (insertions and deletions) we can perform between doublings.
- ▶ What if it's all deletions?
 - ▶ In this case the table becomes progressively less and less full.
 - ▶ Solution: Shrink the table.
- ▶ How, exactly?

Shrinking a Hash Table

Shrinking Tables

- ▶ What should our strategy for reducing the size of the table be?
- ▶ How about “if $n < m/2$ make $m' = m/2$ ”?
- ▶ What if the next operation is an insertion?
 - ▶ Double the table size!
 - ▶ Then a deletion?
 - ▶ Halve the table!
 - ▶ Insertion?
 - ▶ Double...
 - ▶ We now have $\Theta(n)$ operations for each change in the data.
- ▶ Instead use “if $n < m/4$ make $m' = m/2$ ”.



Constant Time?

- ▶ Although, in our example, $T(n)$ is in $\Theta(1)$; for some operations the actual cost is in $\Theta(n)$.
- ▶ What does this mean if we have a real-time application?
 - ▶ Every so often we get an insertion /deletion that takes a really loooooong time.
 - ▶ Can we remedy this so that **every** operation is in $\Theta(n)$?
- ▶ The answer is Yes!
 - ▶ We simply adopt the following strategy...
 - ▶ ...when a table starts to become full—perform the table doubling in the background.
 - ▶ Keep two sets of the data until you either actually need the double size or until the panic is over.

Hashing With Chaining Considered Bad

- ▶ There is still one small issue with this method.
 - ▶ We have a hybrid data structure—an array of linked lists.
- ▶ A second approach uses just a simple array.
- ▶ Clearly, we still have a potential problem with collision—two keys which hash to the same value.
- ▶ We resolve this with a technique known as *Open Addressing*.

Open Addressing

An alternative to chaining

Open Addressing

- ▶ We wish to hash n items into an array with m slots.
- ▶ We may only store one item per slot.
- ▶ Clearly, $m \geq n$.
- ▶ We insert an item into the table using an iterative technique known as *probing*.

Probing

- ▶ This process works as follows: (for insertion)
- ▶ Set hash function to starting value, h_0
repeat
 - calculate $probe = hash(key)$
 - if $table(probe)$ contains data then
 - go to the next hash function
 - else
 - store the item in $table(probe)$
 - fiuntil we have stored the item
- ▶ This means we must have a sequence of hash functions, h_0, h_1, h_2, \dots
- ▶ ... or a hash function which produces a sequence of values.

The Hash Function

- ▶ Our new hash function requires two arguments:
 - ▶ The key;
 - ▶ The iteration count.
- ▶ Thus: **probe**=OpenHash(**key**, **count**)
- ▶ Here:
 - ▶ **key** is a valid element of U , the universe of keys;
 - ▶ **count** is a non-negative integer.
- ▶ As usual, $0 \leq \mathbf{probe} < m-1$.

- ▶ In addition, we want our hash function to have the following property:
- ▶ For any arbitrary key k the sequence of m probes:
 - ▶ $h(k, 0), h(k, 1), h(k, 2), \dots, h(k, m-1)$;
- ▶ Must be a permutation of the integers:
 - ▶ $0, 1, 2, \dots, m-1$.
- ▶ This property guarantees that we must eventually find a vacant slot to insert the item into.
- ▶ Clearly, the sequence of probes must be different for different keys.
- ▶ We can see this with an example.

Example: Insertion with Open Addressing

- Consider the following table:

k	$h(0,k)$	$h(1,k)$	$h(2,k)$	$h(3,k)$	$h(4,k)$	$h(5,k)$	$h(6,k)$	$h(7,k)$	$h(8,k)$	$h(9,k)$
899	9	8	5	6	0	7	8	2	4	1
950	5	7	4	9	2	3	1	6	8	0
12	3	8	7	2	5	9	1	6	0	4
367	7	1	2	3	4	5	6	8	9	0
359	2	1	9	5	6	7	3	8	0	4
980	4	7	1	8	9	3	0	5	2	6
229	0	8	2	7	1	6	3	9	4	5
598	8	6	3	5	0	7	9	1	4	2
838	6	2	6	7	1	3	8	2	0	2
549	9	8	4	6	7	5	0	1	2	3

- Let us insert the keys into our hash table in order

	k	h(0,k)	h(1,k)	h(2,k)	h(3,k)	h(4,k)	h(5,k)	h(6,k)	h(7,k)	h(8,k)	h(9,k)
→	899	9	8	5	6	0	7	8	2	4	1
→	950	5	7	4	9	2	3	1	6	8	0
→	12	3	8	7	2	5	9	1	6	0	4
→	367	7	1	2	3	4	5	6	8	9	0
→	359	2	1	9	5	6	7	3	8	0	4
→	980	3	7	1	8	9	4	0	5	2	6
→	229	0	8	2	7	1	6	3	9	4	5
→	598	8	6	3	5	0	7	9	1	4	2
→	838	6	2	4	7	1	3	8	2	0	2
→	549	9	8	4	6	7	5	0	1	2	3

0	1	2	3	4	5	6	7	8	9
229	980	359	12	549	950	838	367	598	899

Search with Open Addressing

- ▶ The procedure used to search using open addressing is similar to insertion.
- ▶

```
count=0
repeat
    probe=hash(key, count)
    if table(probe)==key then
        return item
    else
        count++
fi
until table(probe)==empty or count==n
return not found
```
- ▶ This is pretty straightforward.

Deletion with Open Addressing

- ▶ When we get to deletion we have a new problem.

- ▶

```
count=0
repeat
    probe=hash(key, count)
    if table(probe)==key then
        delete item
        return
    else
        count++
fi
until table(probe)==empty or count==n
return not found
```

- ▶ How, exactly, do we delete the item?

Deletion...

- ▶ If we simply replace the item with our empty value we will have an issue:
 - ▶ What if the key we next search for is after the probe corresponding to the deleted key's location.
 - ▶ If, in our previous example, we delete 899, where $h(899,0)=9$, and then search for 549, where the sequence of hash values are 9, 8, 4...
 - ▶ We test $D(9)$ and discover it has the value **empty**.
 - ▶ We conclude that 549 is not in the table.
 - ▶ Wrong! It is in $D(4)$.
- ▶ To fix this we need a second special value, **deleted**.

Deletion concluded

- Our deletion process becomes:

```
► count=0
  repeat
    probe=hash(key, count)
    if table(probe)==key then
      table(probe)==deleted
      return
    else
      count++
  fi
until table(probe)==empty or count==n
return not found
```

- This fixes search but introduces a problem with insertion.

Insertion Revisited.

- ▶ We note that we can insert a new item into the dictionary in two circumstances:
 - ▶ `D(i)==empty`
 - ▶ `D(i)==deleted`
- ▶ We modify our insert process as follows:
 - ▶ `count=0`
repeat
 - `probe=hash(key, count)`
if `table(probe)==empty` or `table(probe)==deleted` then
store item in `table(probe)`
return
 - else
count++
 - fi
 - until `count==n`
return no room
- ▶ Now we can insert into the first vacant slot, empty or deleted, that we find in the table.

Search Revisited

- ▶ Because **empty** and **deleted** are different, we do not have to modify our search procedure.
- ▶ The search will skip over deleted records because they do not match the key but will still terminate when it reaches an empty record.

Hash Functions

Open Addressing Hash Functions

- ▶ One question remains.
- ▶ Can we find a function $h(k, i)$ which is:
 - ▶ Easy to compute;
 - ▶ Produces a permutation of $\{0, 1, \dots, m-1\}$ as i varies over $\{0, 1, \dots, m-1\}$?
- ▶ Let us examine two possible strategies.

Strategy I: Linear Probing

- ▶ In this approach we simply take a standard hash function, $h(k)$ and compute the probe $p(k, i)$ as follows:
 - ▶ $p(k, i) = (h(k) + i) \bmod m$
- ▶ In other words, we simply look at sequential entries in the dictionary starting at the entry corresponding to $h(k)$.
 - ▶ This is certainly easy to compute.
 - ▶ It does satisfy the permutation.
- ▶ Is it any good?
 - ▶ No!
 - ▶ It produces sets of consecutive occupied slots.
 - ▶ Clustering.
- ▶ The bigger the cluster, the more likely it is to be hit..
 - ▶ ...and it gets even bigger!

Strategy II: Double Hashing

- ▶ In this strategy we have two standard hash functions, $h_1(k)$ and $h_2(k)$.
- ▶ We compute $p(k)$, our probe value as follows:
 - ▶ $p(k, i) = (h_1(k) + i \times h_2(k)) \bmod m$.
- ▶ Do we still satisfy our requirements?
 - ▶ This is still easy to compute.
 - ▶ Do we always get a permutation?
 - ▶ No.
 - ▶ Unless we are clever in how we define h_2 .

Choosing h_2

- ▶ We need $h_2(k)$ to be relatively prime to m .
- ▶ I.e. $h_2(k)$ and m must have no common factors except 1.
- ▶ This is easy in many cases.
- ▶ If we select m to be a power of 2; say $m = 2^r$ then all we need is for $h_2(k)$ to always be an odd number.
- ▶ For example, if we have a standard hash function $h'(k)$, we can create $h_2(k)$ as follows:
 - ▶ $h_2(k) = (2h'(k) + 1) \bmod m$

Table Doubling.

- ▶ Once again, we need to expand the dictionary whenever it becomes too full.
- ▶ What does “too full” mean in this case?
- ▶ We define the occupancy of a table, α , to be the ratio of n , the number of entries to m , the number of slots.
 - ▶ $\alpha = n/m$
 - ▶ $0 \leq \alpha \leq 1$
- ▶ We can show that the average cost of an operation on a table with occupancy α is in $\Theta(1/(1 - \alpha))$.
- ▶ In practice we want this value to be reasonably close to 1 so we double as soon as α exceeds 0.5 or thereabouts.
- ▶ This keeps operations between $\Theta(1)$ and $\Theta(2)$.

An Important Note on α

- ▶ When calculating the occupancy value, α , we must count slots with a value of **deleted** as containing data.
- ▶ This is because some operations, notably searching, treat deleted records as still containing data.
- ▶ Slots containing **deleted** may be removed in two ways:
 - ▶ Being overwritten with valid data as a result of an insert operation;
 - ▶ Being cleaned up when the table is expanded.
- ▶ If we did not count **deleted** records in calculating α we could have a notionally empty table in which every slot was **deleted**.
- ▶ Search (and delete) in this table would be $\Theta(m)$, not $\Theta(1)$, as we might expect.

Chaining vs. Open Addressing

- ▶ So, which is the better scheme?
- ▶ Open Addressing:
 - ▶ Uses less memory—no need for pointers;
 - ▶ Is faster—provided α is kept below 0.5;
 - ▶ Is a little harder to implement and understand.
 - ▶ Is clean—one data structure, the array.
- ▶ Chaining:
 - ▶ Uses more memory;
 - ▶ Is faster—if we are not careful with open addressing.
 - ▶ Is a little easier to implement and understand.
 - ▶ Is a bit messy—arrays of linked lists.
- ▶ I know where my vote goes!