

CSCI203

Week 6 – Lecture B

Looking for Text (In all the right places)

- ▶ Consider the problem of *String Searching*:
 - ▶ Given a text, t , is the subtext, s , present in it?
- ▶ This problem occurs in many real-life applications:
 - ▶ grep;
 - ▶ find in a text editor;
 - ▶ Genome matching;
 - ▶ Google search.
- ▶ There are a wide number of techniques to achieve this.
- ▶ Let us look at a couple of examples.

Linear Search

The Naïve Approach: Linear Search

- ▶ The simplest possible approach is linear search:
 - ▶ Try to match s starting at each location in t .
 - ▶ for i in $0 \dots \text{length}(t) - \text{length}(s)$
 - $j=0$
 - while $j < \text{length}(s)$ do
 - if $(s[j] \neq t[i+j])$ break
 - $j++$
 - od
 - if $j == \text{length}(s)$ print(" string found starting at location " i)
 - rof
- ▶ We can see this with an example.

Linear Search: an Example

- ▶ Let t be the string “harry happened to have a hard hand”.
- ▶ Let s be the string “hard”.
- ▶ The search proceeds as follows:

h a r r y h a p p e n e d t o h a v e a h a r d h a n d

h a r r y h a p p e n e d t o h a v e a h a r d h a n d

h a r r y h a p p e n e d t o h a v e a h a r d h a n d

h a r r y h a p p e n e d t o h a v e a h a r d h a n d

...

h a r r y h a p p e n e d t o h a v e a h a r d h a n d

h a r r y h a p p e n e d t o h a v e a h a r d h a n d

Linear Search \neq Linear Time Search

- ▶ The outer loop in our algorithm is repeated $|t| - |s|$ times.
 - ▶ Typically the string t is much longer than the string s , so this is $\Theta(t)$.
- ▶ The inner loop is repeated up to $|s|$ times for each time round the outer loop.
 - ▶ This is $\Theta(s)$
- ▶ The total number of comparisons is $\Theta(|s| \times |t|)$.
- ▶ Is this the best we can do?
- ▶ The best we can possibly do is $\Theta(|s| + |t|)$;
 - ▶ we have to at least look at each string!
- ▶ Can we actually achieve this goal of a linear time algorithm?

Linear Time Search

Linear Time Search

- ▶ To do this we will use hashing.
- ▶ We compare the hash of string s with the hash of each substring of t with the same length:
 - ▶

```
hash_s = hash(s)
for i in 0...length(t)-length(s)
    hash_t = hash(t[i..i+length(s)-1])
    if hash_s == hash_t then
        brute-force compare s and the substring
        if they match print(" string found starting at
                                location ", i)
    fi
rof
```


Linear Time Search

- ▶ This algorithm takes linear time, provided:
 - ▶ The hash function only collides rarely;
 - ▶ The hash function takes constant time to compute;
 - ▶ Independent of the length of string s !
- ▶ Surely, the second requirement is impossible.
- ▶ To hash a string of length $|s|$ must take $\Theta(|s|)$ operations.
 - ▶ Yes?
 - ▶ No!
- ▶ Not if we are clever.

Clever Hashing

- ▶ We note that we need $\Theta(|s|)$ time to compute $h(s)$.
- ▶ We also need $\Theta(|s|)$ time to compute $h(t[0..|s|-1])$, the initial substring of t .
- ▶ The trick is to compute the hash of each successive substring of t in constant time.
- ▶ If we look closely at these substrings, we see an interesting feature:
 - ▶ Successive substrings differ only by two characters.
- ▶ The first character of the first substring;

h	a	r	r
---	---	---	---
- ▶ The last character of the next substring.

a	r	r	y
---	---	---	---

Rolling Hash

- ▶ Maybe we can define a hash function which, given $h(\text{"harr"})$ can compute $h(\text{"arry"})$ in constant time.
- ▶ Let us define a *rolling hash* function, $r()$, so that:
 - ▶ $h(\text{"arry"}) = r(h(\text{"harr"}), \text{"h"}, \text{"y"})$
 - ▶ We compute the hash of the next substring by removing the first and appending the new last characters;
 - ▶ In this case we remove "h" and append "y".
- ▶ If we can compute a rolling hash in constant time than we can do string matching in linear time.
- ▶ How?

Karp-Rabin String Search

- ▶ The Karp-Rabin algorithm looks like this:

```
▶ hash_s=hash(s)
  hash_t=hash(t[0..length(s)-1])
  for i in 0...length(t)-length(s)
    if hash_s == hash_t then
      brute-force compare s and the substring
      if they match print(" string found starting at
                           location " i)
    fi
    hash_t=roll(hash_t,t[i],t[i+length(s)])
  rof
```

- ▶ The function $\text{roll}(h,p,s)$ computes the rolling hash of the next substring given the hash of the existing substring, h , with the prefix p removed and the suffix, s , appended.
- ▶ We need only find a suitable function $\text{roll}()$.

How We Roll

- ▶ One popular way to compute `roll()` is to use something called the Rabin fingerprint.
- ▶ We start by treating each symbol in the alphabet as an integer – use the ASCII code for example.
- ▶ We then find a random prime number $>$ the size of the alphabet—let's pick 257.
- ▶ We now compute $h(\text{"harr"})$ as:
 - ▶ $257^3 \cdot 104 + 257^2 \cdot 97 + 257^1 \cdot 114 + 257^0 \cdot 114$
 - ▶ $= 1,771,793,837$
 - ▶ Note: "h" = 104, "a" = 97 and "r" = 114.

The Next Hash

- ▶ Ok, so given that $h(\text{"harr"}) = 1,771,793,837$ how do we get $h(\text{"arry"})$?
- ▶ It's easy:
- ▶ Simply compute $r(h,p,s) = 257 \cdot (h - 257^3 \cdot p) + s$
- ▶ $257 \cdot (1,771,793,837 - 257^3 \cdot 104) + 121$
- ▶ In this case the result is 1,654,094,526 which is exactly the same as $h(\text{"arry"})$
- ▶ $257^3 \cdot 97 + 257^2 \cdot 114 + 257^1 \cdot 114 + 257^0 \cdot 121$
- ▶ Note: if these values become too large, we can reduce them modulo m , where m is a convenient value—say 2^{15} or 2^{31} .

Efficient?

- ▶ We can compute our hash values for s and the initial substring of t using *compact evaluation*.
- ▶ $p^{k-1}.c_1 + p^{k-2}.c_2 + \dots + p.c_{k-1} + c_k$
- ▶ This requires a lot of multiplication!
- ▶ It can be re-written as...
- ▶ $h = c_k + p(c_{k-1} + p(c_{k-2} + \dots + p(c_3 + p(c_2 + pc_1)) \dots))$
- ▶ Where $k = |s|$ and c_i is the i^{th} character of s .
- ▶ This requires $|s| - 1$ multiplications and $|s| - 1$ additions.

Efficiency!

- ▶ If we precompute $q=p^{k-1}$ we can find the next hash value, h' as:
- ▶ $h'=p.(h-q.c_i)+c_j$ where we remove character i and add character j .
- ▶ This requires only 1 multiplication and 1 addition.
- ▶ Constant time.
- ▶ Thus we have $\Theta(|s|)$ operations to perform the initial hashes and $|t|-|s| * \Theta(1)$ operations to do the rehashing.
- ▶ Overall: our algorithm operates in $\Theta(|s|+|t|)$ time.
- ▶ We win!