

Ian Piper
CSCI203


Algorithms and Data Structures

Week 8 – Lecture A

Graphs

Graphs

- ▶ A graph consists of a set, V , of points, called *Vertices* or *Nodes*, and a set, E , of *Edges*, also called *Arcs*, each of which contains a pair of vertices.
- ▶ $G=(V, E)$
 - ▶ $V=\{v_1, v_2, v_3, v_4, v_5, v_6, v_7\}$
 - ▶ $E=\{\{v_1, v_2\}, \{v_1, v_4\}, \{v_2, v_3\}, \{v_4, v_5\}, \{v_4, v_7\}, \{v_5, v_6\}, \{v_5, v_7\}, \{v_6, v_7\}\}$
 - ▶ This is an *undirected* graph
 - ▶ $E=\{(v_1, v_2), (v_1, v_4), (v_2, v_3), (v_4, v_1), \{v_4, v_5\}, (v_5, v_7), (v_6, v_5), (v_7, v_4), (v_7, v_5)\}$
 - ▶ This is a *directed* graph.




Representing a Graph



- ▶ The best way to represent a graph in a computer depends on how the graph is to be used.
- ▶ The obvious representation—an array of vertices and an array of edges—is probably the worst possible way!
- ▶ Two common representations are:
 - ▶ Adjacency List;
 - ▶ Adjacency Matrix.

Adjacency Lists

- ▶ An adjacency list, L , is an array (or hash table), of length $|V|$, of linked lists where:
 - ▶ The list stored in L_i consists of all the vertices directly reachable from vertex i .
- ▶ The graph shown to the right...
- ▶ ...has the following adjacency list representation:



Adjacency List




Adjacency Matrix

- ▶ An adjacency matrix is a $|V| \times |V|$ array containing zeros and ones.
 - ▶ Note: $|V|$ is the number of vertices in V .
 - ▶ A zero is stored in location (i, j) if there is no edge between v_i and v_j .
 - ▶ A one is stored in location (i, j) if there is an edge between v_i and v_j .
- ▶ Note:
 - ▶ If G is an nondirected graph, the array is symmetric.
 - ▶ $a(i, j) = a(j, i)$.

Adjacency Matrix

0	1	0	1	0	0	0
1	0	1	0	0	0	0
0	1	0	0	0	0	0
1	0	0	0	1	0	1
0	0	0	1	0	1	1
0	0	0	0	1	0	1
0	0	0	1	1	1	0



Abstract Notation

- ▶ If we ignore the detail of how, exactly, we store the graph we can represent it using the following abstract notation.
- ▶ $G = \{\text{Adj}(v_1), \text{Adj}(v_2), \dots, \text{Adj}(v_{|V|})\}$
- ▶ Here, G is defined as the addressable set of elements $\text{Adj}(v_i)$ where $\text{Adj}(v_i)$ is the set of vertices directly reachable from vertex v_i .
- ▶ This representation allows us to use any appropriate data structure to implement the graph;
 - ▶ Arrays;
 - ▶ Hash tables;
 - ▶ Matrices.

Graph Search

Graph Search

- ▶ A common problem involving graphs is *Graph Search*.
- ▶ This involves 'exploring' the graph in some systematic way starting at vertex s and visiting the other reachable vertices by following edges from one vertex to the next.
- ▶ This can be done in more than one way, as we shall see.
- ▶ Graph search has many real life applications, including:
 - ▶ Web crawling;
 - ▶ Network broadcast;
 - ▶ Social networking;
 - ▶ Garbage collection;
 - ▶ Solving puzzles and games.

Breadth-First Search (BFS)


- ▶ Our goal is to list all the vertices that are reachable from some starting node $s \in V$ by following edges.
- ▶ To do this in $\Theta(|V| + |E|)$ time.
- ▶ Strategy:
 - ▶ List all the nodes reachable from s in 0 moves;
 - ▶ List all the new nodes reachable from s in 1 move;
 - ▶ List all the new nodes reachable from s in 2 moves;
 - ▶ ...
 - ▶ List all the new nodes reachable from s in n moves.
- ▶ Note: a new node is one that has not already been visited.
- ▶ This avoids duplicate nodes in our result.

BFS Algorithm

```
► Procedure BFS(s, adj):  
  v: set of vertex  
  q: queue of vertex  
  v={}  
  q.enqueue(s)  
  while q is not empty:  
    current = q.dequeue()  
    add current to v  
    for each n in adj(current)  
      if n is not in v then  
        q.enqueue(n)  
      fi  
    rof  
  elihw  
  return v  
end BFS
```

BFS, an Example

► BFS(v_1 , Adj)



► Current node: \mathbf{c}

v_6

► Visited nodes: \mathbf{v}

v_1 v_2 v_4 v_3 v_5 v_7 v_6

► Pending nodes: \mathbf{q}

v_1 v_2 v_4 v_3 v_5 v_7 v_6

Keeping Track

- ▶ Sometimes we wish to keep track of how we got to each node in a graph.
- ▶ To achieve this we need only make a small change to the BFS algorithm.
 - ▶ We add an array, or hash table, P , indexed by each vertex we reached, containing the *parent* of each vertex, the vertex from which we reached it.
 - ▶ Clearly:
 - ▶ The parent of our starting vertex s is empty; $p(s) = \text{null}$
 - ▶ The parent of each vertex directly connected to s is s .
 - ▶ The modified BFS is shown on the next slide.

BFS with Parents


```
► Procedure BFS_Parent(s, adj):  
  v: set of vertex  
  q: queue of vertex  
  p[vertex]: array of vertex  
  v={}  
  q.enqueue(s)  
  p[s]=null  
  while q is not empty:  
    current = q.dequeue()  
    add current to v  
    for each n in adj(current)  
      if n is not in v then  
        q.enqueue(n)  
        p[n]=current  
      fi  
    rof  
  elihw  
  return v  
end BFS_Parent
```


Using the Parents

- ▶ To find the path from vertex s to any given vertex, d , we simply list the sequence:
 - ▶ d ,
 - ▶ $p[d]$
 - ▶ $p[p[d]]$
 - ▶ ...
 - ▶ $p[p[...p[[d]]...]]$
 - ▶ Until we reach s .
- ▶ The reverse of this list is the path we seek.

BFS with Parent Tracking, an Example

► BFS(v_1 , Adj)



► Current node: **c**

v_6

► Visited nodes: **v**

v_1 v_2 v_4 v_3 v_5 v_7 v_6

► Pending nodes: **q**

v_1 v_2 v_4 v_3 v_5 v_7 v_6

► Parent node: **p**

ϕ v_1 v_1 v_2 v_4 v_4 v_5

Efficiency of BFS

- ▶ We note that, in the worst case, we visit all vertices in the graph.
- ▶ For each vertex in the graph we traverse each of its edges.
- ▶ The complexity of BFS is, therefore, $O(|V| + |E|)$.
 - ▶ Strictly speaking, for undirected graphs, $O(|V| + 2 \times |E|)$
 - ▶ For directed graphs, $O(|V| + |E|)$.

Depth First Search (DFS)


- ▶ The goal of DFS is the same as BFS:
 - ▶ List all vertices reachable from the starting vertex s .
- ▶ Strategy:
 - ▶ For each vertex v in $\text{Adj}(s)$;
 - ▶ Perform DFS on v .
 - ▶ This is a recursive implementation of DFS.
- ▶ We can also perform DFS using a non-recursive algorithm.
- ▶ All we need to do is replace the queue from BFS with a stack.

DFS Algorithm

```
► Procedure DFS(s, adj):  
  v: set of vertex  
  k: stack of vertex  
  p[vertex]: array of vertex  
  v={}  
  k.push(s)  
  p[s]=null  
  while k is not empty:  
    current = k.pop()  
    add current to v  
    for each n in adj(current)  
      if n is not in v then  
        k.push(n)  
        p[n]=current  
      fi  
    rof  
  elihw  
  return v  
end DFS
```

DFS, an Example

► DFS(v_1 , Adj)



► Current node: **c**

v_3

► Visited nodes: **v**

v_1 v_4 v_7 v_6 v_5 v_2 v_3

► Pending nodes: **k**


v_3 v_5 v_6

Some Notes

- ▶ BFS and DFS will visit the same vertices given the same starting vertex.
 - ▶ All that changes is the order they are visited.
- ▶ The edges traversed in performing a DFS or BFS form a tree.
 - ▶ If the graph is connected this is a *spanning* tree.
- ▶ BFS and DFS work on directed as well as on undirected graphs.

Traversing a Maze


- ▶ If we modify DFS to track parents we can use it as a means of finding our way through a maze.
- ▶ Consider the following maze:
- ▶ We can represent it as a graph with 18 vertices, a, b, \dots, p, s and x .
- ▶ An edge exists in this graph if two vertices are adjacent and do not have a wall between them.
- ▶ We start the maze at vertex s and exit to vertex x .



A(maze)d

Traversing the Maze

- ▶ If we try directions in the sequence:
 - ▶ Down;
 - ▶ Up;
 - ▶ Left;
 - ▶ Right.
 - ▶ (we need to stack in the reverse order; R, L, U, D)
- ▶ The DFS proceeds as follows:



When DFS Fails

- ▶ If the graph is not connected, or is directed and not strongly connected, we may not reach every vertex with a single call to DFS.
- ▶ In this case we will need to call DFS repeatedly, once for each vertex we have not yet visited.
- ▶ We can do this easily with the following procedure, DFS_All:


DFS_All

```
► Procedure DFS_All(g, adj)
    v: set of vertex
    v={}
    for each s in g
        if s not in v then
            DFS(s, adj)
        fi
    rof
end DFS_All
```

► Note: we must remove the line $\mathbf{v}=\{\}$ from procedure DFS.

DFS_All, an Example

- Consider the following graph:




- If we run DFS_All in alphabetical order we get the following result:

- v:

a	b	e	d	c	f
---	---	---	---	---	---
- p:

ϕ	a	b	e	ϕ	c
--------	---	---	---	--------	---



So, Why Both?

- ▶ If both BFS and DFS can be used to search a graph why bother with both of them.
- ▶ The answer is simply that, although they both perform graph search, they provide different information about the graph.
- ▶ For example, BFS will always find the *shortest* path from s to x , assuming a path exists, while DFS will not.
 - ▶ Shortest is the fewest vertices in the path.
- ▶ Many applications of DFS rely on the way it can be used to classify edges.


Edge Classification

Edge Classification

- ▶ If we perform a DFS on a graph we can classify the edges of a graph:
 - ▶ Tree edges: these form part of the search tree (or forest);
 - ▶ Forward edges: these lead from a vertex to a descendant;
 - ▶ Backward edges: these lead from a vertex to an ancestor;
 - ▶ Cross edges: these are all the edges that are left—they connect unrelated vertices.

Edge Classification: an Example

- ▶ In the Graph shown to the right;
- ▶ The edges are classified as follows:
- ▶ Tree edges (in red);
- ▶ Forward edges (in green);
- ▶ Backward edges (in orange);
- ▶ Cross edges (in cyan).



How to Tell

- ▶ How do we determine the type of an edge in our DFS procedure?
- ▶ The parent structure, $p[]$, gives us the tree edges.
 - ▶ If $p[i] = j$ then the edge from vertex j to vertex i is a tree edge.
- ▶ If we encounter a vertex that is currently in the recursion stack, k , then the corresponding edge is a backward edge.

How to Tell

- ▶ To identify forward and cross edges we need to record *when* we visit a vertex—the sequence in which the vertices are pushed onto the vertex stack, k .
- ▶ With this sequence number we can determine whether a remaining edge is a forward or a cross edge:
 - ▶ If the sequence number of the vertex we are visiting is larger than the sequence number of the current vertex this is a forward edge;
 - ▶ Otherwise this is a cross edge.
- ▶ We could also record the finishing time of the vertices—the sequence in which the vertices are popped from the vertex stack, k .

DFS Algorithm

```
► Procedure DFS_Timed(s, adj):
    in: static integer = 0
    out: static integer = 0
    v: set of vertex
    k: stack of vertex
    p[vertex]: array of vertex
    k.push(s)
    s.start=in++
    p[s]=null
    while k is not empty:
        current = k.pop()
        s.end=out++
        add current to v
        for each n in adj(current)
            if n is not in v then
                k.push(n)
                n.start=in++
                p[n]=current
            fi
        rof
    elihw
    return v
end DFS_Timed
```



Why?

- ▶ Why is edge classification important?
- ▶ We will consider two problems in graph theory:
 - ▶ Cycle detection;
 - ▶ Topological sorting.
- ▶ Each of these depends on edge classification for its solution.

Cycle Detection

Cycle Detection

- ▶ We define a cycle as any sequence of edges that form a closed loop in a graph.
- ▶ E.g. in the graph to the right...
- ▶ The edges (b, e) , (e, d) and (d, b) form a cycle.
- ▶ Finding these cycles is trivial once we have classified all the edges:
 - ▶ G has a cycle if, and only if, $\text{DFS}(G, \text{Adj})$ contains at least one backward edge.



Finding Cycles

- ▶ Once we find a backward edge, (v_0, v_k) , finding the cycle it forms part of is easy:
 - ▶ Simply follow the parent sequence, back from v_0 until we reach v_k .
 - ▶ These edges plus the backward edge must form the cycle.
- ▶ Our next problem, *Topological Sort*, requires that the graph be acyclic:
 - ▶ It has no cycles.
 - ▶ We now know how to do this.
- ▶ We will look at topological sort by way of a simple application, *job scheduling*.

Topological Sort

Job Scheduling.

- ▶ Consider a directed, acyclic graph (DAG) in which:
 - ▶ The vertices each represent a task to be performed;
 - ▶ The edges represent an order in which the tasks may be performed.
 - ▶ If edge (v_i, v_j) exists in the graph, task v_i must be completed before we start task v_j .
- ▶ Our problem is to find a *feasible sequence* of tasks:
 - ▶ An order in which the tasks may be performed which does not conflict with the order required by all of the edges.
- ▶ You can only do one task at a time.
- ▶ Consider the task of getting dressed:

Getting Dressed: Vertices


- ▶ We can label our tasks (vertices), in alphabetical order:
 - ▶ G: put on glasses;
 - ▶ H: put on shirt;
 - ▶ J: put on jacket;
 - ▶ O: put on shoes;
 - ▶ S: put on socks;
 - ▶ T: put on trousers.
 - ▶ U: put on underpants;
 - ▶ V: put on vest;
 - ▶ W: put on watch;
- ▶ We can then determine the edges:

Getting Dressed: Edges

- ▶ We can define our edges as follows:
 - ▶ (V, H): vest before shirt;
 - ▶ (U, T): underpants before trousers;
 - ▶ (S, O): socks before shoes;
 - ▶ (T, O): trousers before shoes;
 - ▶ (H, J): shirt before jacket.
- ▶ Note vertices G and W do not appear in the edge list:
 - ▶ We can do these tasks at any time.
- ▶ Taken together, the vertices and edges form a DAG.

Getting Dressed: the Graph


- ▶ The resulting graph is shown to the right:
- ▶ We perform our topological sort as follows:
 - ▶ Conduct a DFS of the graph;
 - ▶ List the vertices in reverse order of finishing time.
- ▶ In practice we construct a list of the vertices as each is popped from the stack and then print the list backwards.



Why is it So?

- ▶ We can prove that the procedure we describe produces a correct ordering:
- ▶ Theorem: if G contains the edge (u, v) then vertex v will be popped before vertex u .
- ▶ Proof: consider two cases.
 - ▶ Case 1: u starts before v .
 - ▶ In this case we must visit v before we are finished with u .
 - ▶ Because edge (u, v) exists, vertex v will be visited while vertex u is still on the stack
 - ▶ Case 2: u starts after v .
 - ▶ In this case we cannot visit v while u is still on the stack.
 - ▶ Because the graph is acyclic, there is no path from v to u .

Topological Sort: an Example



- ▶ Given the “getting dressed” graph:
- ▶ Our DFS (performed in alphabetical order) goes like this:
 - ▶ Push G: no edges—pop G.
 - ▶ Push H: edge (H, J)
 - ▶ Push J: no edges—pop J
 - ▶ Back at H: no more edges—pop H
 - ▶ Push O: no edges—pop O
 - ▶ Push S: no unvisited vertices—pop S
 - ▶ Push T: no unvisited vertices—pop T
 - ▶ Push U: no unvisited vertices—pop U
 - ▶ Push V: no unvisited vertices—pop V
 - ▶ Push W: no edges—pop W
- ▶ Our topological sort is:
W, V, U, T, S, O, H, J, G.

G J H O S T U V W

One of Many.

- ▶ The example we just saw gives us one of many possible topological sorts for the “getting dressed” graph.
- ▶ If we visited the vertices in a different order we would probably get a different sort order.
- ▶ E.g. reverse alphabetical order results in the topological sort G, S, U, T, O, V, H, J, W.
- ▶ Verify:
 - ▶ This is really what you get;
 - ▶ This is a valid sequence of operations.