

CSCI203

Week 9 – Lecture B

Tweaking Dijkstra

- ▶ As we have already seen, Dijkstra's algorithm provides an effective solution strategy for solving the single source/all destinations version of the shortest path problem.
- ▶ We will now look at a few simple modification of Dijkstra that will:
 - ▶ Improve its practical performance;
 - ▶ and/or extend its range of applicability.

Overall Efficiency

- ▶ The algorithm as you have seen it so far has efficiency $O(|V|^2 + |E|)$.
- ▶ But I said it was $O(|V| \cdot \log |V| + |E|)$.
- ▶ How come?
- ▶ The answer is simple:
 - ▶ As presented we find the next vertex to select by searching a list of candidate vertices and finding the vertex with minimum D value.
 - ▶ This is a linear search process; $O(|V|)$.
 - ▶ We do this for each vertex, also $O(|V|)$.
 - ▶ This is $O(V^2)$.
- ▶ So, how do we improve on this?

Reaching Peak Efficiency

- ▶ The answer is surprisingly simple.
- ▶ Replace the candidate list/array C with...
- ▶ ...a priority queue (aka a heap), ordered on $D(v)$.
- ▶ Now:
 - ▶ Finding the best candidate is $O(1)$;
 - ▶ Updating C is $O(\log |V|)$.
- ▶ Now, over all vertices, we have $|V| \cdot \log |V|$.

Single Destination.

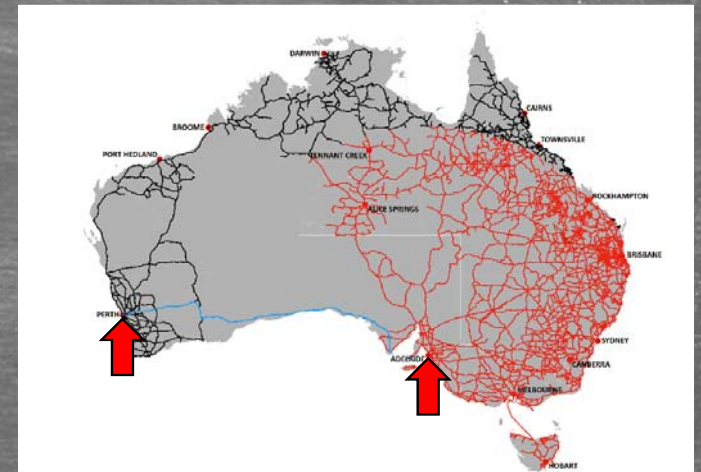
- ▶ What do we do if, rather than looking for the shortest paths from a start vertex, s , to all other vertices, we wish to find the shortest path from s to a specific goal vertex, g ?
- ▶ The answer is easy:
 - ▶ Stop when vertex g becomes a member of S , the selected set.
- ▶ This means that we do not waste time with any vertex further away from s than g .
- ▶ This usually reduces the total running time of the algorithm.
 - ▶ Why *usually*?

All Sources—Single Destination

- ▶ In this case, rather than finding paths from a starting vertex, s , to all other reachable vertices we are looking for the shortest paths to some goal vertex, g , from all possible starting vertices.
- ▶ How do we do this?
 - ▶ Run Dijkstra backwards.
- ▶ Specifically:
 - ▶ Redefine $\text{Adj}(v)$ to be the list of set of vertices leading to vertex v ...
 - ▶ ...instead of reachable from v ;
 - ▶ Start with $D(g)=0$...
 - ▶ ...instead of $D(s)=0$;
 - ▶ Let $P(v)$ indicate the next vertex in the path...
 - ▶ ...instead of the prior vertex;
 - ▶ Let the selected set, S , start at $\{g\}$...
 - ▶ ...instead of $\{s\}$.

The Problem with Dijkstra

- ▶ There is a big problem with using Dijkstra on the single source/single destination problem:
 - ▶ The order in which the vertices are added.
- ▶ Consider the following graph:
 - ▶ Say we want to get from Adelaide...
 - ▶ ...to Perth.
- ▶ Dijkstra will add all of the closer vertices first:
- ▶ Before we ever get close to the path we seek.



Fixing the Problem

- ▶ So, how do we remedy this?
- ▶ Before we get to the answer let us take a step back.
- ▶ Let us generalize Dijkstra's algorithm.
- ▶ The key step in the algorithm is on the process by which we select the next vertex.
 - ▶ Specifically, we select the vertex in the candidate set, C , for which the overall distance to the vertex from the source vertex, s , is minimized.
 - ▶ $D(v) = P(s, v)$.
- ▶ Note that this does not involve the goal vertex, g .

Eyes on the Goal

- ▶ What if we could bias the selection towards the goal in some way.
- ▶ We can!
- ▶ Essentially, we select the minimum not simply of $P(s, v)$ but, instead of $P(s, v) + H(v, g)$.
- ▶ This new function, H , is a *heuristic*; an estimate of the remaining distance from each candidate vertex, v , to the goal vertex, g .
- ▶ What is a good estimator?

A Good Heuristic

- ▶ We require $H(v, g)$ to have one key property:
 - ▶ $H(v, g) \leq P(v, g)$;
 - ▶ The heuristic estimate must never exceed the actual shortest path length.
 - ▶ This requirement guarantees that the final path we find is still the correct answer.
- ▶ In our example we have a ready-made heuristic...
- ▶ ...The Euclidean (straight-line) distance between v and g .
- ▶ Provided the graph behaves according the rules of geometry there can never be a shorter path than this.

A*

- ▶ The heuristic modification to the vertex selection rule changes Dijkstra's algorithm into an example of what is called the A* algorithm.
- ▶ Although, in the worst case, A* is no faster than Dijkstra in practice it will generally represent a substantial improvement.
- ▶ Note: the trick to A* is finding a good heuristic.
- ▶ The nearer that $H(v, g)$, the estimated minimum path length from v to g , is to $P(v, g)$, the actual minimum path length, the faster A* will find the solution.

Through the Looking Glass.

- ▶ Because of the order in which we saw them, it is easy to think of A^* as a generalization of Dijkstra's algorithm.
- ▶ This is not the only, or perhaps even the best, way to view this.
- ▶ Consider instead this viewpoint:
- ▶ Dijkstra's algorithm is simply a special case of A^* :
- ▶ The one with the worst possible choice of H .
- ▶ Specifically, $H(v, g)=0$.